

Shreya Reddy

510-646-1363 | shreyared3@gmail.com | [linkedin.com/in/shreyar3](https://www.linkedin.com/in/shreyar3)

SUMMARY

Senior AI/ML Engineer with 4+ years building production ML systems across healthcare, fintech, and telecom. Specializes in end-to-end MLOps, real-time feature engineering, LLM-powered RAG applications, and multi-model ensemble deployment on AWS SageMaker and Google Cloud Vertex AI. Track record of 30–50% efficiency gains and multi-million-dollar revenue impact through scalable ML infrastructure, GenAI agents, and regulatory-compliant model governance.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, Bash, Scala

ML & AI: XGBoost, LightGBM, Gradient Boosting, LSTM, PyTorch, HuggingFace Transformers, Scikit-learn, LangChain, LangGraph, SHAP, Pinecone, FAISS, OpenAI API, RAG, Fine-tuning

Cloud & Platforms: AWS SageMaker, Google Cloud Vertex AI, BigQuery, Snowflake, Databricks, Azure ML

Data Engineering: Apache Kafka, Apache Spark, Airflow, AWS Glue, Amazon Redshift, Pub/Sub, S3, dbt

MLOps & Deployment: MLflow, Evidently AI, Docker, Kubernetes, FastAPI, Cloud Run, Apigee, CI/CD, Git, Terraform

EXPERIENCE

Cardinal Health

Dublin, OH

AI/ML Engineer

Jun 2024 – Present

- Migrated pharmaceutical supply chain ML pipelines from Databricks to Google Cloud Vertex AI, rebuilding drug shortage risk scoring using XGBoost and LSTM ensemble for scalable production deployment.
- Designed real-time feature pipelines on Pub/Sub and Kafka ingesting EHR events, supplier fill-rate streams, and external signals into BigQuery, replacing overnight batch processing with sub 3 minute feature latency.
- Built an internal RAG based procurement assistant using LangChain and Pinecone, enabling supply chain teams to query drug shortage policies and inventory reports in natural language.
- Deployed risk score endpoints via Cloud Run through an Apigee-secured REST API with automated retraining on Vertex AI Pipelines, consumed by IntelLogix dashboards and SAP ERP systems for auto reorder triggers.

Discover Financial Services

Riverwoods, IL

Machine Learning Engineer

May 2024 – Jun 2024

- Developed a three-model credit risk scoring system on AWS SageMaker for card application decisioning, limit optimization, and collections propensity using XGBoost, LightGBM, and Gradient Boosting.
- Engineered a centralized feature store on SageMaker Feature Store with signals including payment ratio, utilization trend, and delinquency history across batch and real time scoring workloads.
- Integrated SHAP explainability via SageMaker Clarify to generate adverse action reasons for declined applications, meeting ECOA and Regulation B compliance requirements.
- Maintained automated retraining and drift monitoring using SageMaker Pipelines and MLflow across weekly production deployments.

Tech Mahindra

Hyderabad, India

Data Engineer

Sep 2020 – Jul 2023

- Built batch ingestion pipelines using AWS Glue and PySpark to migrate telecom customer usage and billing data from on-prem Oracle systems into Amazon S3.
- Designed Redshift data warehouse tables and optimized SQL transformation scripts to support incremental loads, enabling business intelligence teams to run churn and revenue analytics.
- Automated daily pipeline scheduling using Apache Airflow across S3, Glue, and Redshift, handling failure alerts and retry logic to maintain data freshness for client reporting SLAs.

EDUCATION

George Mason University

Fairfax, VA

Master of Science in Data Analytics Engineering

Aug 2023 – May 2025

Jawaharlal Nehru Technological University

Hyderabad, India

Bachelor of Technology in Information Technology

Aug 2019 – Jun 2023