

Shreya Reddy

510-646-1363 | shreyared3@gmail.com | [LinkedIn](#) | [GitHub](#)

SUMMARY

AI/ML Engineer with 4+ years developing, shipping, and supporting RAG-based agents and LLM-powered applications for non-technical enterprise users across healthcare and fintech. Specializes in prompt engineering, document ingestion pipelines, chunking strategy, answer quality validation, and cross-functional AI enablement using OpenAI API, LangChain, and LlamaIndex. Track record of translating ambiguous operational problems into production-ready AI tools with measurable adoption impact.

EXPERIENCE

Cardinal Health

Dublin, OH

AI/ML Engineer

Jun 2025 – Present

- Built and maintained RAG-based AI agents using OpenAI API (GPT-4), LangChain, LangGraph, and Pinecone on AWS, implementing document ingestion pipelines, chunking strategies, vector embeddings, and prompt template libraries, cutting manual policy lookup time by 65% for supply chain teams.
- Validated answer quality and tuned multi-agent behavior through structured testing of retrieval accuracy, prompt engineering configurations, and chunking parameters, iterating on agent design to improve response reliability for non-technical end users.
- Designed and delivered AI enablement programs including onboarding guides, user documentation, and training materials for 150+ cross-functional stakeholders across procurement, operations, and finance.
- Deployed AI agent endpoints on AWS with REST APIs and automated retraining pipelines, tracking RAG pipeline performance and adoption metrics across enterprise integrations.

Discover Financial Services

Riverwoods, IL

Machine Learning Engineer

May 2024 – Jun 2025

- Developed a three-model payment decisioning system on AWS SageMaker for card application scoring, limit optimization, and collections propensity using XGBoost, LightGBM, and Gradient Boosting, processing 2M+ monthly applications and improving credit approval accuracy by 18%.
- Wrote clear use-case documentation for model outputs and SHAP-based adverse action reasons, translating complex ML decisions into plain language for non-technical Risk and Compliance teams under ECOA and Regulation B.
- Maintained model monitoring, drift detection, and automated retraining using MLflow and SageMaker Pipelines with CI/CD via GitHub Actions, tracking performance and adoption metrics across weekly production deployments.

Tech Mahindra

Hyderabad, India

Data Engineer

Sep 2020 – Jul 2023

- Built ETL pipelines using AWS Glue and PySpark to migrate 800GB+ of telecom customer usage and billing data from on-prem Oracle systems into Amazon S3, supporting downstream analytics for business intelligence teams.
- Automated pipeline scheduling using Apache Airflow across S3, Glue, and Redshift, handling failure alerts and retry logic to achieve 99.7% SLA compliance for client reporting.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, Bash

AI Agents & LLMs: OpenAI API (GPT-4), LangChain, LangGraph, LlamaIndex, RAG, multi-agent orchestration, prompt engineering, document ingestion, chunking strategies, vector embeddings, answer quality validation, Pinecone, FAISS, prompt template libraries

Cloud & Platforms: AWS, AWS SageMaker, Amazon S3, GCP, BigQuery, Databricks

Data Engineering: Apache Kafka, Apache Spark, AWS Glue, Apache Airflow, Amazon Redshift

MLOps & Tools: MLflow, Evidently AI, CI/CD, GitHub Actions, Git, Docker, FastAPI

EDUCATION

George Mason University

Fairfax, VA

Master of Science in Data Analytics Engineering

Aug 2023 – May 2025

Jawaharlal Nehru Technological University

Hyderabad, India

Bachelor of Technology in Information Technology

Aug 2019 – Jun 2023